

Transform Data

dplyr

Your Turn 1

Alter the code to select just the `n` column:

```
select(babynames, n)
```

```
# A tibble: 1,924,665 x 1
      n
  <int>
1  7065
2  2604
3  2003
4  1939
5  1746
6  1578
7  1472
8  1414
9  1320
10 1288
# ... with 1,924,655 more rows
```

Quiz

Which of these is NOT a way to select the `name` and `n` columns together?

```
select(babynames, -c(year, sex, prop))
```

```
# A tibble: 1,924,665 x 2
  name      n
  <chr>  <int>
1 Mary    7065
2 Anna    2604
3 Emma    2003
4 Elizabeth 1939
5 Minnie   1746
6 Margaret 1578
7 Ida      1472
8 Alice    1414
9 Bertha   1320
10 Sarah   1288
# ... with 1,924,655 more rows
```

```
select(babynames, name:n)
```

```
# A tibble: 1,924,665 x 2
  name      n
  <chr>  <int>
```

```

1 Mary      7065
2 Anna      2604
3 Emma      2003
4 Elizabeth 1939
5 Minnie    1746
6 Margaret  1578
7 Ida       1472
8 Alice     1414
9 Bertha    1320
10 Sarah    1288
# ... with 1,924,655 more rows

```

```
select(babynames, starts_with("n"))
```

```
# A tibble: 1,924,665 x 2
```

```

  name      n
  <chr>   <int>
1 Mary    7065
2 Anna    2604
3 Emma    2003
4 Elizabeth 1939
5 Minnie  1746
6 Margaret 1578
7 Ida     1472
8 Alice   1414
9 Bertha  1320
10 Sarah  1288
# ... with 1,924,655 more rows

```

```
select(babynames, ends_with("n"))
```

```
# A tibble: 1,924,665 x 1
```

```

  n
  <int>
1  7065
2  2604
3  2003
4  1939
5  1746
6  1578
7  1472
8  1414
9  1320
10 1288
# ... with 1,924,655 more rows

```

Your Turn 2

Use `filter`, `babynames`, and the logical operators to find:

- All of the names where `prop` is greater than or equal to 0.08
- All of the children named “Sea”

```
filter(babynames, prop >= 0.08)
```

```
# A tibble: 3 x 5
```

	year	sex	name	n	prop
	<dbl>	<chr>	<chr>	<int>	<dbl>
1	1880	M	John	9655	0.0815
2	1880	M	William	9532	0.0805
3	1881	M	John	8769	0.0810

```
filter(babynames, name == "Sea")
```

```
# A tibble: 4 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr> <int>    <dbl>
1  1982 F    Sea      5 0.00000276
2  1985 M    Sea      6 0.00000312
3  1986 M    Sea      5 0.0000026
4  1998 F    Sea      5 0.00000258
```

Your Turn 3

Use Boolean operators to return only the rows that contain:

- *Boys* named Sue
- Names that were used by exactly 5 or 6 children in 1880
- Names that are one of Acura, Lexus, or Yugo

```
filter(babynames, sex == 'M', name == "Sue")
```

```
# A tibble: 52 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr> <int>    <dbl>
1  1917 M    Sue      7 0.0000073
2  1927 M    Sue      5 0.0000043
3  1928 M    Sue      5 0.00000438
4  1930 M    Sue      5 0.00000443
5  1931 M    Sue      6 0.00000561
6  1932 M    Sue      7 0.00000652
7  1933 M    Sue      7 0.00000686
8  1934 M    Sue     14 0.0000132
9  1935 M    Sue     13 0.0000122
10 1936 M    Sue      9 0.00000846
```

... with 42 more rows

```
filter(babynames, year == 1880, (n == 5) | (n == 6))
```

```
# A tibble: 455 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr> <int>    <dbl>
1  1880 F    Abby      6 0.0000615
2  1880 F   Aileen      6 0.0000615
3  1880 F    Alba      6 0.0000615
4  1880 F   Alda      6 0.0000615
5  1880 F   Alla      6 0.0000615
6  1880 F  Alverta      6 0.0000615
7  1880 F    Ara      6 0.0000615
8  1880 F  Ardelia      6 0.0000615
9  1880 F  Ardella      6 0.0000615
```

```
10 1880 F      Arrie      6 0.0000615
# ... with 445 more rows
```

```
filter(babynames, (name == "Acura") | (name == "Lexus") | (name == "Yugo"))
```

```
# A tibble: 57 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr> <int>    <dbl>
1 1990 F    Lexus    36 0.0000175
2 1990 M    Lexus    12 0.00000558
3 1991 F    Lexus   102 0.0000502
4 1991 M    Lexus    16 0.00000755
5 1992 F    Lexus   193 0.0000963
6 1992 M    Lexus    25 0.0000119
7 1993 F    Lexus   285 0.000145
8 1993 M    Lexus    30 0.0000145
9 1994 F    Lexus   381 0.000195
10 1994 F    Acura     6 0.00000308
# ... with 47 more rows
```

```
filter(babynames, name %in% c("Acura", "Lexus", "Yugo")) # Same as above
```

```
# A tibble: 57 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr> <int>    <dbl>
1 1990 F    Lexus    36 0.0000175
2 1990 M    Lexus    12 0.00000558
3 1991 F    Lexus   102 0.0000502
4 1991 M    Lexus    16 0.00000755
5 1992 F    Lexus   193 0.0000963
6 1992 M    Lexus    25 0.0000119
7 1993 F    Lexus   285 0.000145
8 1993 M    Lexus    30 0.0000145
9 1994 F    Lexus   381 0.000195
10 1994 F    Acura     6 0.00000308
# ... with 47 more rows
```

Help Me

What is the smallest value of n? What is the largest?

```
arrange(babynames, n) # Smallest
```

```
# A tibble: 1,924,665 x 5
  year sex  name      n      prop
  <dbl> <chr> <chr>    <int>    <dbl>
1 1880 F    Adelle     5 0.0000512
2 1880 F    Adina      5 0.0000512
3 1880 F    Adrienne   5 0.0000512
4 1880 F    Albertine  5 0.0000512
5 1880 F    Alys       5 0.0000512
6 1880 F    Ana        5 0.0000512
7 1880 F    Araminta   5 0.0000512
8 1880 F    Arthur     5 0.0000512
9 1880 F    Birtha     5 0.0000512
10 1880 F    Bulah      5 0.0000512
```

```
# ... with 1,924,655 more rows
arrange(babynames, desc(n))

# A tibble: 1,924,665 x 5
  year sex   name      n  prop
  <dbl> <chr> <chr>   <int> <dbl>
1  1947 F     Linda  99686 0.0548
2  1948 F     Linda  96209 0.0552
3  1947 M     James  94756 0.0510
4  1957 M    Michael 92695 0.0424
5  1947 M    Robert  91642 0.0493
6  1949 F     Linda  91016 0.0518
7  1956 M    Michael 90620 0.0423
8  1958 M    Michael 90520 0.0420
9  1948 M     James  88588 0.0497
10 1954 M    Michael 88514 0.0428
# ... with 1,924,655 more rows
```

Your Turn 4

Use `%>%` to write a sequence of functions that:

1. Filters `babynames` to just the girls that were born in 2017, *then...*
2. Selects the `name` and `n` columns, *then...*
3. Arranges the results so that the most popular names are near the top.


```
babynames %>%
  filter(sex == 'F', year == 2017) %>%
  select(name, n) %>%
  arrange(desc(n)) -> girls_2017
```

Your Turn 5

1. Trim `babynames` to just the rows that contain **your name** and **your sex**
2. Trim the result to just the columns that will appear in your graph (not strictly necessary, but useful practice)
3. Plot the results as a line graph with `year` on the x axis and `prop` on the y axis

```
babynames %>%
  filter(name=="Derek", sex=="M") %>%
  select(year, prop) %>%
  ggplot() +
  geom_line(aes(x=year, y=prop), color = "darkgreen") +
  labs(title = "Where's Derek? (Popularity of the name Derek)",
       x = "Year",
       y = "Proportion of Males Born") +
  theme_minimal()
```

Where's Derek? (Popularity of the name Derek)



Take aways

- Extract variables with `select()`
- Extract cases with `filter()`
- Arrange cases, with `arrange()`